Jibram Jimenez-Loza

Software Engineer, ML & AI Systems

San Francisco, CA | JibramSWE@gmail.com | jibram.net | linkedin.com/in/jibram

SUMMARY

Software engineer with 5+ years building ML and AI systems end to end, from models and data pipelines to the product and release infrastructure around them. At Google, built ML reserve-pricing models that generated $10M+ in revenue, plus the MLOps pipelines that trained and tuned them. Currently building release infrastructure for Letta's desktop app.

EXPERIENCE

Nob Hill Labs (Contracting for Letta, an AI research lab) **Aug 2026 - Present**
Senior Software Engineer *Remote (SF Bay Area)*

- Proposed and built a packaged end-to-end (e2e) test stage in GitHub Actions for Letta's Electron desktop app, testing the fully built app users install and gating production promotion on failure.
- Caught a couple of packaging regressions (asar, native modules) that the existing e2e suite missed in its first week, including an open user-reported bug; fixed both.

Chelys / HumaniCore **Mar 2026 - Present**
Independent Software Engineer *Remote (SF Bay Area)*

- Built and launched the marketing sites and user intake funnels for two pre-seed startups (chelys.io, humanicore.ai), setting up DNS, GitHub-connected Vercel preview and production deploys, and initial AWS/Supabase services.
- Built the server-side intake flow in Astro with input scrubbing and rate limiting; advised founders on product direction and the feasibility and scaling of services.

Datavant **Sep 2025 - Dec 2025**
Software Engineer *Remote (SF Bay Area)*

- Extended an ELT pipeline (Python, SQL, Snowflake) to support a client-specific schema, including database migrations and intermediate-table transformations, taking the client live on a production searchable database.
- Integrated the new data into the existing customer-matching and heuristic-ranking workflow across Python/FastAPI services, AWS RDS, and Snowflake.

Google **Oct 2021 - Nov 2024**
Software Engineer *Remote*

- Engineered ML reserve-pricing models for sparse travel-ads auctions and deployed region-specific variants, generating 2.5%+ revenue lift ($10M+).
- Extended large-scale MLOps pipelines for deep neural networks (TensorFlow, TFX, distributed C++ data processing with Flume) to new ad verticals, tuning model capacity, layers, and training stability for each.
- Built a simulation pipeline on Google Vizier (Bayesian optimization) that automated hyperparameter retuning after model or data changes, replacing manual engineer tuning with configs optimized for maximum profit and minimal click loss.

Ancestry **Jan 2020 - Oct 2021**
Associate Software Engineer *Remote*

- Improved checkout conversion by using funnel drop-off and error-rate dashboards to find struggling user groups, then shipping and A/B testing React component fixes.
- Built go-to-market features, including multi-region and multi-language validation, that opened new international markets and contributed to 5%+ user growth.

SKILLS

ML & Data: TensorFlow, TFX, hyperparameter tuning (Vizier), Snowflake, SQL
Languages: Python, C++, TypeScript, JavaScript
Cloud & Delivery: AWS, GCP, GitHub Actions, CI/CD, Vercel, Supabase
Backend & Apps: FastAPI, Node.js, gRPC, React, Electron, Astro, Jest, Cypress

EDUCATION

University of California, Merced **2019**
B.S. in Computer Science and Engineering | GPA: 3.6 | Dean's Honors

Yonsei University, Seoul **2017 - 2018**
Study Abroad | Coursework: Artificial Intelligence, Computer Vision